

Giving the leaves back to the tree: Using random forest models for imputation

Damon C. Roberts*

Abstract

Other fields such as the Biomedical Sciences and Computer Science use and advocate for Random Forest models as a flexible tool for imputing sparse datasets. To my knowledge, the use of this tool has yet to enjoy wide usage in political science. Implementations of Random Forest models in the multiple imputation paradigm referred to as MICE (or as it is sometimes called fully conditional specification) are flexible and loosen the distributional assumption required for popular techniques used by political scientists. Additionally, implementation of these models is relatively easy. The MICE implementation in Python, R, and STATA include Random Forest models for the use of imputation. In this manuscript, I use simulated and real world survey data to examine the conditions under which Random Forest models perform well. I compare its performance to that of AMELIA and of other MICE models.

*Ph.D Student, Department of Political Science, University of Colorado Boulder, UCB 333, Boulder, CO 80309-0333. Damon.Roberts-1@colorado.edu.

Acknowledgements: I would like to thank Andrew Q. Philips for our many conversations and his advice for this manuscript. I would also like to thank Andrew Baker for encouraging me to take the time and the space to write this manuscript.

Replication Code can be found at: <https://github.com/DamonCharlesRoberts/imputation-with-random-forests>

1 Introduction

Missing values are common in social science data. King and colleagues (King et al. 2001) estimate that political scientists lose about one-third of their data in their complete case regression analyses due to missing data. There are many reasons as to why missing values arise in our data. In surveys, these are referred to item-nonresponse and pose threats to unbiased estimates to public opinion research when researchers utilize listwise deletion (LWD) in their regression analyses (Weisberg 2005); particularly when the researcher is able to predict the cause of the missingness with an observed cause - which the data are referred to as "missing at random" (King et al. 2001).¹ More generally, missing values come about in other types of data as a result of the unit's (e.g. country, respondent, politician) attempt to obfuscate information, data collector error, or researcher error. When common and not from a stochastic data generating process (DGP), missing data pose threats to causal inference either through model identification with non-full rank matrices or sample representativeness.

Unsurprisingly a number of political scientists generate tools and learn from other fields to deal with these challenges. Many of the most popular approaches to imputation carry assumptions that scholars have to satisfy. There are four primary approaches that political scientists use in dealing with missing data: (1) listwise deletion or complete case analysis (LWD), (2) simple mean or median based imputation, (3) hot deck or regression based imputation, or (3) multiple imputation (MI). In this manuscript, I advocate for the application of flexible machine learning approaches to imputing missing values using a variant of the MI framework called Multiple Imputation through Chained Equations (MICE). Those who use machine learning models for predictive

¹Which is a common occurrence, more so than missingness caused entirely by random chance - which, if this condition is met, gives unbiased estimates with LWD.

regressions often use random forest models due to their flexibility as a result of their fully non-parametric nature. Fields like the biomedical sciences treat missing values as out-of-sample predictions that random forest models predict.

Not only do other fields use random forest models for imputation, but it is relatively easy to implement as well with contemporary computational and software capabilities. Using the Multiple Imputation by Chained Equations (MICE) paradigm, there are a number of single-line implementations in common data analysis software such as STATA, R, and Python (see van Buuren and Groothuis-Oudshoorn 2011).

This manuscript compares the utility of random forest models for imputation to other common imputation tools used in political science. The next section provide more details on the common imputation approaches that political scientists currently use. Following that, the following section describes random forest models and their utility for predicting missing data. I then move into applied examples where I compare the performance of these random forest models to other common imputation approaches on simulated data and an applied example with the 2017-2020 panel of the World Values Survey. I then close with a discussion of recommendations for when one should use the reigning popular techniques or the random forest application.

2 Common Approaches to dealing with missing data, imputation, and MI in Political Science

2.1 MCAR, MAR, and MNAR

Missing data arise in different forms. There are three key ways researchers describe missing data. The first form that missing data take are Missing Completely At Random

(MCAR). This means that the data generating process for the missingness is random - there are no observed or unobserved causes of missingness. The second form that missing data take are Missing At Random (MAR). In MAR, these data are missing due to some observed cause. Some argue that MAR is much more common given the state of how large most contemporary social science data sets are (Schunk 2008). The third form that missing data take are Missing Not At Random (MNAR).² MNAR happens when the cause for missingness is explained both by observed and unobserved causes. Given that missing data take different forms, researchers use a few different approaches to deal with these data.

2.2 Dealing with missingness

At the time of writing King and colleagues (2001) estimated that 94% of political scientists use LWD to deal with missing data. In short, LWD does not seek to impute missing values. Instead, if the DV or any of the covariates in a regression model for a given observation are missing, the researcher does not include that observation in the analysis. Traditionally, scholars argue that LWD performs best (in terms of reducing the resulting bias in the researcher's subsequent regression models) when the data are MCAR. If the data are MAR or MNAR, deleting observations with missing data are generating bias in one's regression estimates through a form of unit non-response bias (see Weisberg 2005).³ A meta-analysis of comparative and international political economy papers that use LWD demonstrates that political scientists have much, upwards of 50%, more Type I error than we would expect as a result of how we implement LWD (Lall 2016). Others

²Sometimes called Non-ignorable (NI).

³Many others make a similar argument, but with different, non-survey, terminology (see King et al. 2001; Schunk 2008; Azur et al. 2011, to name a few).

push against this claim and instead argue that LWD does not inherently generate bias for non-MCAR data, but that researchers neglect to control for the cause of MAR or MNAR (Arel-Bundock and Pelc 2018).

Like LWD, simple imputation techniques like mean- and median-based imputation do not reduce the chances of biased regression estimates. These approaches sometimes called interpolation and extrapolation are commonly used with panel data. If you have missing data for an observation in one panel, you can take the same observations' response in a previous panel and a latter panel. You then take the mean or the median of that particular observation for that variable. Other approaches seek to reduce this MAR-based bias through conditioning on other variables.

Hot Deck approaches to imputing missingness are regression-based in that they define the dependent variable as the one the researcher is attempting to impute and use variables thought to predict the cause of missingness in MAR contexts (Schunk 2008). You often use a minimal number of variables to condition on in this approach. In many cases, the precise mechanism generating missingness is often quite difficult to triangulate. As a result, if you fail to provide the right model specification when the data are MAR, you still often end up with biased regression estimates anyhow.

MI seeks to solve this issue through using the entire dataset for imputing missing values (Rubin 1996). This approach uses the other variables in the dataset to generate a joint posterior distribution of all possible missing values for that particular observation. Many assume that most social science data sets are sufficiently large enough to condition on the mechanism generating MAR (Schunk 2008). Unlike the other approaches, MI also generates uncertainty around the imputed values (Rubin 1996) which enables the researcher to be more transparent about the validity of those imputed values and to include that uncertainty in the researcher's subsequent statistical analyses (King et al.

2001; von Hippel 2015). A very popular implementation of MI in political science is Honaker, King, and Blackwell's (2011) AMELIA II software (Lall 2016). This useful tool provides a computationally fast and simple process for imputation. Compared to the other approaches to missing data, AMELIA II performs quite well (2011; 2014). MI, however, often requires a set of distributional assumptions for the joint distribution - often the multivariate normal (Honaker, King and Blackwell 2011).

There is a variant to MI which seeks more computational efficiency and loosens some of the distributional assumptions required. This variant is called Multiple Imputation through Chained Equations (MICE). MICE performs quite well for large imputation tasks. MI struggles to impute values when there is missingness in the other variables of the dataset (Kropko et al. 2014) - which is quite common. This is because the way that the joint distribution is constructed. If there are missing values in the variables to be included in the joint distribution, this introduces bias in that distribution. MICE tries to get around this problem in a few steps as described by Azur and colleagues (Azur et al. 2011). First, it performs a simple imputation for every missing value in the entire dataset. These are the placeholder values. The second step involves identifying one variable to impute. Once this has been done, it then removes those placeholder values. The third step then involves regressing the observed values of the variable on the other variables in the model and replace the predicted values generated from the regression model for the missing values. The fourth step is to then repeat steps two and three for each variable in the data set with missing values - this constitutes a single iteration. As a fifth step, you perform between five and ten iterations.⁴

The regression models that one may use in MICE are as numerous as the regres-

⁴Though, the exact number of recommended iterations used in MICE are still up for debate (see van Buuren and Groothuis-Oudshoorn 2011; ?).

sion models a researcher may choose from when engaged in a statistical analysis. This means, however, that the assumptions and the performance of the model one uses for the imputation are the same as they are in standard statistical analyses. Selection of the model to use in MICE should be carefully thought out. One valuable model for generating out-of-sample predictions is a form of ensemble machine learning model called Random Forests.

3 The utility of random forest models for imputing political science data

Random forest models are concerned with calculating a fixed out of sample prediction. This is in line with the argument that multiple imputation should not be evaluated on the model's correctness but on the model's ability to predict a fixed (true) value (Rubin 1996). Here, I think of missing data as the out of sample predictions that are intended to be estimated. With random forest models, you train the model on a training data set (often randomly generated through cross validation) and then fit the model on the testing set (Hastie, Tibshirani and Friedman 2009). OLS models often perform relatively poorly on making out-of-sample predictions as they are BLUE assuming the data on hand are relatively representative of the population. As we have MAR data, this assumption likely fails and any out-of-sample predictions are likely to be biased. OLS may also generate out-of-bounds predictions for non-continuous data (Long 1997; Gelman, Hill and Vehtari 2021).

There are other models which provide within-bounds predictions such as logistic models, however, as generalized linear models, they still assume a linear functional

form. Random forest models both provide within-bounds predictions and they are fully non-parametric (Hastie, Tibshirani and Friedman 2009); meaning they do not assume a functional form and consequently a joint distribution.

On other performance metrics, random forest models, as an ensemble method, provide much more accurate predictions than single models, such as CART (Montgomery and Olivella 2018). Rather than generating the single estimate from a single model, ensemble models like random forests, calculate multiple models and learn from their performance. Specifically, random forest models take bootstrapped samples of the dataset and fit multiple trees (a single tree model being one which successively fit models that refine a trade-off between the bias and efficiency).

In the context of MICE, political scientists can use random forest models to make accurate predictions with fewer assumptions and leniency in terms of conditioning on the cause of MAR in the data set. Recall that within each MICE iteration, one performs a model predicting (imputing) a missing value based on the other variables in the dataset. Explicitly, this means that you are running a model per variable with missing data. Given that random forests are non-parametric, within each MICE iteration, the relationship between the variable to be imputed and those used to make the predictions can be non-linear and can take many different distributional forms. This is a significant advancement on traditional MI which assumes a joint normal distribution. Additionally, this is an advancement on other MICE models which may assume linear relationships. Furthermore, using random forests in MICE has the advantage over hot deck procedures for those unsure about the precise DGP for MAR in a variable. Figure 1 provides a brief summary of using random forests to impute missing values with MICE.

Single MICE iteration

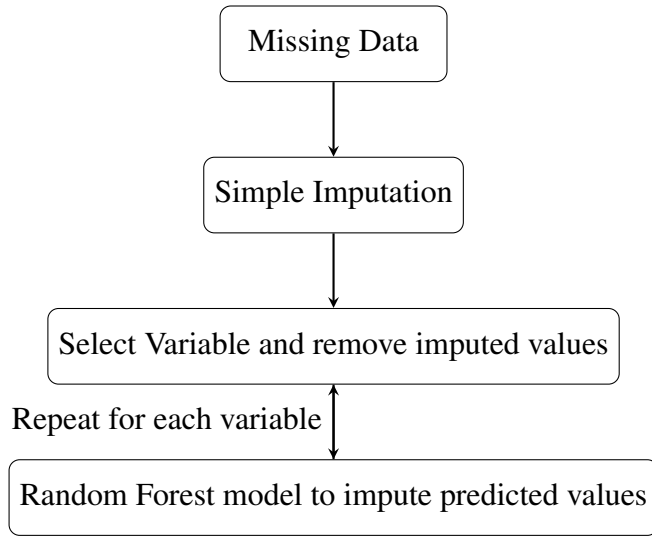


Figure 1: MICE with Random Forest models

For example, say we have a 3x3 matrix, m , where:

$$m = \begin{bmatrix} NA & 1 & 2 \\ 1 & NA & 2 \\ 1 & 2 & NA \end{bmatrix}$$

The first step would be to do a simple imputation that returned:

$$m = \begin{bmatrix} \mathbf{1.5} & 1 & 2 \\ 1 & \mathbf{1.5} & 2 \\ 1 & 2 & \mathbf{1.5} \end{bmatrix}$$

The second step would be to remove m_{11} yielding:

$$m = \begin{bmatrix} NA & 1 & 2 \\ 1 & \mathbf{1.5} & 2 \\ 1 & 2 & \mathbf{1.5} \end{bmatrix}$$

And then running a random forest model to predict the value of m_{11} .

This second step yields:

$$m = \begin{bmatrix} 1.4 & 1 & 2 \\ 1 & \mathbf{1.5} & 2 \\ 1 & 2 & \mathbf{1.5} \end{bmatrix}$$

This is then repeated for the other columns in the matrix. Once completed for each column in the matrix, this constitutes 1 iteration. This is then done between five and 10 more times. This then gives you a distribution of imputed values for each cell that were originally missing.

In the next section, I illustrate the use of random forest models for political scientists by demonstrating a simulated and real-world application of a random forest implementation of MICE. The next section also compares this implementation's performance to that of other common approaches to handling missing data in political science.

4 Application of MICE with random forests

4.1 Simulated Data

1. Distributional Performance (MVN versus other distributions)
2. MCAR, MAR, MNAR Performance

3. Variable structure (i.e. continuous, ordinal, nominal, discrete)??

4.2 In the wild: World Values Survey 2017-2020 Panel

5 Conclusion

The imputation model used here is a form of multiple imputation. As a result, it does have its limitations. As King and colleagues (2001) outline, listwise deletion preforms better than MI when:

- (1) The analysis model is conditional on X ..., and the functional form is known to be correctly specified ...
- (2) There is NI [MNAR] missingness in X , so that [the algorithm] can give incorrect answers, and no Z variables are available that could be used in an imputation stage to fix the problem.
- (3) Missingness in X is not a function of Y , and unobserved omitted variables that affect Y do not exist...
- (4) The number of observations left after listwise deletion should be so large that the efficiently loss from listwise deletion does not counterbalance the biases induced by the other conditions.

Random forest models not for multiple imputation are especially prone to a decline in performance when X is sparse. Given that this already is a problem for MI models, this is certainly a significant drawback to the implementation of random forests for MI. If one has a sparse data set, using random forest models for MI should certainly be reconsidered.

References

- Arel-Bundock, Vincent and Krzysztof J. Pelc. 2018. "When Can Multiple Imputation Improve Regression Estimates?" *Political Analysis* 26(2):240–245.
- Azur, Melissa J., Elizabeth A. Stuart, Constantine Frangakis and Philip J. Leaf. 2011. "Multiple imputation by chained equations: what is it and how does it work?" *International Journal of Methods in Psychiatric Research* 20(1):40–49.
- Gelman, Andrew, Jennifer Hill and Aki Vehtari. 2021. *Regression and Other Stories*. New York: Cambridge University Press.
- Hastie, Trevor, Robert Tibshirani and Jerome Friedman. 2009. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer Series in Statistics 2nd ed. New York: Springer.
- Honaker, J, G King and M Blackwell. 2011. "Amelia II: A program for missing data, R package version 1.5., 2012." *Journal of Statistical Software* 45(7):1–3.
- King, Gary, James Honaker, Anne Joseph and Kenneth Scheve. 2001. "Analyzing Incomplete Political Science Data: An Alternative Algorithm for Multiple Imputation." *American Political Science Review* 95(1).
- Kropko, Jonathan, Ben Goodrich, Andrew Gelman and Jennifer Hill. 2014. "Multiple imputation for continuous and categorical data: Comparing joint multivariate normal and conditional approaches." *Political Analysis* 22(4):497–519.
- Lall, Ranjit. 2016. "How Multiple Imputation Makes a Difference." *Political Analysis* 24(4):414–433.

- Long, Scott J. 1997. *Regression Models for Categorical and Limited Dependent Variables*. Advanced Quantitative Techniques in the Social Sciences Series Thousand Oaks, CA: Sage Publications.
- Montgomery, Jacob M and Santiago Olivella. 2018. “Tree-Based Models for Political Science Data.” *American Journal of Political Science* 62(3):729–744.
- Rubin, Donald B. 1996. “Multiple Imputation After 18+ Years.” *Journal of the American Statistical Association* 91(434):473–489.
- Schunk, Daniel. 2008. “A Markov chain Monte Carlo algorithm for multiple imputation in large surveys.” *AStA* 92(1):101–114.
- van Buuren, Stef and Karin Groothuis-Oudshoorn. 2011. “mice: Multivariate Imputation by Chained Equations in R.” *Journal of Statistical Software* 45(9).
- von Hippel, Paul T. 2015. “New Confidence Intervals and Bias Comparisons Show That Maximum Likelihood Can Beat Multiple Imputation in Small Samples.” *Structural Equation Modeling: A Multidisciplinary Journal* 23(3):422–437.
- Weisberg, Herbert F. 2005. *The Total Survey Error Approach: A Guide to the New Science of Survey Research*. Chicago, IL: Chicago University Press.